



UNIVERSITÀ
di CAMERINO

SuMo

Mutation Testing for Solidity Contracts



Writing Better Tests with SuMo

1. **SUT: Running Example:**
 - Introduction to the SUT;
2. **Code Coverage: Solidity-Coverage**
 - What it tells us about our tests;
 - Why it doesn't tell the whole story.
3. **Mutation Testing: SuMo:**
 - Introduction to SuMo;
 - Running Mutation Testing;
 - Analyzing Live Mutants.



Demo Repository



morenabarboni/sumo-demo

SuMo-Demo

In the repository you will find:

- CampusCoin Setup
- Test Code Examples
- Slides

Prerequisites

Ensure the following are installed:

- Node.js
- npm (comes with Node.js)

SUT: Business Logic

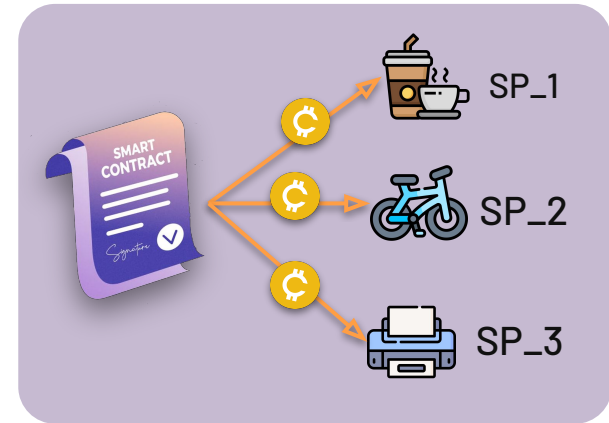
contracts/CampusCoin.sol: A custom **ERC-20 token** that can be used to pay for services around campus (e.g., buying food, borrowing books).

Active Roles:

- **Admin**: Registers Users and Mints new CC;
- **Student**: Pays Service Providers using CC;

Passive Roles:

- **ServiceProvider**: Receives CC payments;
- **University**: Receives a 1% fee on payments.



SUT: Test Environment

`test/CampusCoin.js`: The HardHat test file for CampusCoin.sol


1) Testing Framework

Provides tools and domain-specific utilities to write, organize and execute test cases.



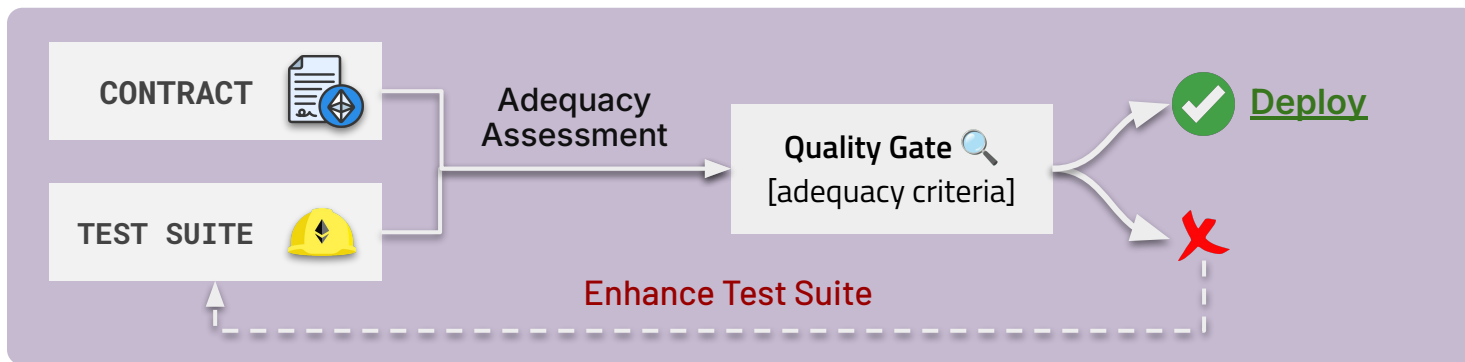
2) Chain Simulator

Creates a **local blockchain environment** to deploy and test contracts without costs.



Goal

Does the CampusCoin.js test suite gives us **sufficient confidence** in the **correctness** of the CampusCoin.sol Smart Contract?



- If the TS meets **adequacy criteria**, the Contract passes the QG;
- Different adequacy criteria = different deployment decisions.

Quality Gate Based on Coverage Criteria

To ensure **high confidence** in the **correctness** of our contract, we define **strict coverage thresholds** as part of our deployment quality gate:



Statement Coverage == 100%

Every statement in the contract must be executed at least once.



Branch Coverage == 100%

Every branch (e.g., **require**) must be evaluated in both directions.

For most contracts, this is feasible and cheaper than the cost of failure.

Coverage Analysis with HardHat


solidity-coverage

chat on [gitter](#) npm@latest v0.8.16  FAILED  codecov 97%  Hardhat  Plugin

Code coverage for Solidity testing

```
22 modifier onlyColonyOwners {  
23   17× if (!this.userIsInRole(msg.sender, 0)) { throw; }  
24   17×  
25 }
```

- For more details about what this is, how it works and potential limitations, see [the accompanying article](#).
- `solidity-coverage` is [Solcover](#)



👉 Let's try it out: **`npx hardhat coverage`**


Coverage Analysis: Results

Our test suite meets the adequacy criteria, looks like we're done...



File	% Stmts	% Branch	% Funcs	% Lines
contracts\ CampusCoin.sol	100	100	100	100
All files	100	100	100	100

Are we really done though?



<https://arstechnica.com> › 2021/12 ▼ Traduci questa pagina

Really stupid “smart contract” bug let hackers steal \$31 million ...

1 dic 2021 — By using the same token for both tokenIn and tokenOut, the hacker greatly inflated the price of the MONO token because the updating of the ...

<https://medium.com> › swlh › the-... ▼ Traduci questa pagina

The Story of the DAO — Its History and Consequences - Medium

In the first few hours of the **attack**, 3.6 million ETH were **stolen**, ... In this exploit, the attacker was able to “ask” the **smart contract (DAO)** to give the ...

<https://www.cnbc.com> › accidenta... ▼ Traduci questa pagina

'Accidental' bug froze \$280 million worth of ether in Parity wallet

8 nov 2017 — Millions of dollars' worth of ether could be frozen on **Parity's** cryptocurrency **wallet** because one individual “accidentally” triggered a **bug**.

Coverage Analysis: Results

File	% Stmts	% Branch	% Funcs	% Lines	Uncovered Lines
contracts\ CampusCoin.sol	100 100	100 100	100 100	100 100	
All files	100	100	100	100	


🛑 **Most Developers are inclined to stop here** 🛑

- **It's a clear and familiar goal** → Intuitive and easily quantifiable.
- **It feels like "enough"** → Reaching 100% feels like a natural "done" signal.
- **It's widely supported** → Well-integrated into most tools.

Coverage Analysis: Why it's not enough

1. **High coverage can give a false sense of security:**
 - It only shows that certain parts of the code were executed;
 - Not that they were meaningfully tested or properly verified.
2. **Coverage metrics are easy to game:**
 - Everyone can write tests that hit every line and branch;
 - We don't know if they're enforcing the right expectations.
 - It is possible to get **100% coverage** with **meaningless tests!**

Example: “Should mint tokens to a student”



CampusCoin implements a simple `mint(address, amount)` function:

- The Admin **mints** new CC tokens to a Student.
- The Student's **token balance** is updated accordingly.

Example: "Should mint tokens to a student"

```
it("Should mint tokens to student", async function () {  
    await campusCoin.mint(student1.address, "100");  
    const studentBalance = await campusCoin.balanceOf(student1.address);  
    expect(studentBalance).to.equal("100");  
});
```

CampusCoin.js: A **test method** verifies the correct behavior of `mint()`.

1. **Simulate Tx:** Admin mints 100 new tokens to a Student;
2. **Assert:** Confirm the correctness of the program behavior:
 - Actual output: student balance
 - Expected output: 100 tokens

Example: "Should mint tokens to a student"

```
it("Should mint tokens to student", async function () {  
    await campusCoin.mint(student1.address, "100");  
    const studentBalance = await campusCoin.balanceOf(student1.address);  
    //expect(studentBalance).to.equal("100");  
});
```

What happens if we remove the assertion?

- We are no longer enforcing any expectation;
- We expect test adequacy to decrease!



Let's try it out: `npx hardhat coverage`

Is Test Coverage Useless?

- **Good** for identifying under-tested parts of the system.
- **Bad** if used as a quality target!


— Inozemtseva and Holmes (2014)



Mutation Testing


Mutation testing offers a stronger alternative, as it evaluates a test suite based on its ability to detect *small faults*.

- Tells us whether **assertions** are meaningful;
- If a **mutant survives**, the test suite **lacks a specific check** that can detect the fault;
- Encourages meaningful, bug-revealing tests.



SuMo - SOLidity MUtator

A flexible and **domain-aware** framework that comprehensively model **Solidity-specific faults** and guides the derivation of meaningful tests.

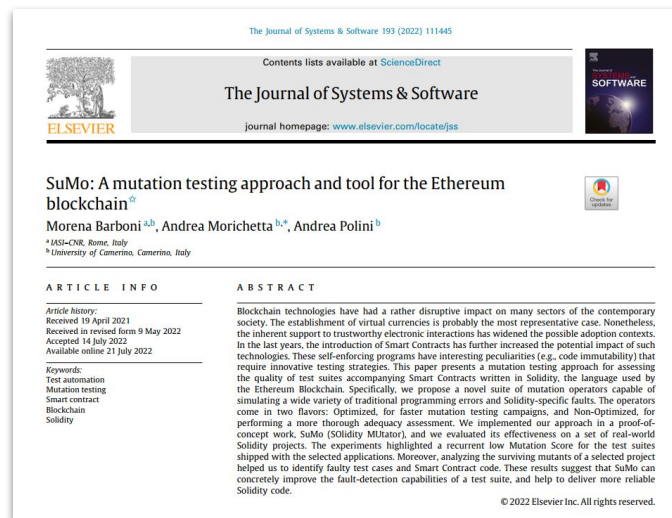


Automated Assessment
For any Solidity project
regardless of frameworks.



40 Mutation Operators


- SOTA tools (PIT);
- Solidity constructs;
- Bugs and pitfalls.



Empirical Evaluation - Key Insights

Test Suites Achieved Low Quality Ratings

Average MS across projects = ~64%



Testers often overlook Solidity-specific constructs:

Traditional Mutants (MS = 68,2 %) | Solidity Mutants (MS = 61,2 %)


Target	MS	Potential Impact
Event	34%	Monitoring of contract behavior
Modifiers	37,7%	Access Control and reusable logic
Exception Handling	40,1%	Management of critical transaction reverting scenarios.
Blockchain Variable	64,5%	Any logic dependant on global blockchain properties (e.g., time)

SuMo can highlight such gaps in Test Suites that would otherwise remain overlooked.

Empirical Evaluation - Key Insights

Test Suites Achieved Low Quality Ratings

Average MS across projects = ~64%




Testers often overlook Solidity-specific constructs:

Traditional Mutants (MS = 68,2 %) | Solidity Mutants (MS = 61,2 %)

Target	MS	Potential Impact
Event	34%	Monitoring of contract behavior
Modifiers	37,7%	Access Control and reusable logic
Exception Handling	40,1%	Management of critical transaction reverting scenarios.
Blockchain Variable	64,5%	Any logic dependant on global blockchain properties (e.g., time)

SuMo can highlight such gaps in Test Suites that would otherwise remain overlooked.

Installing SuMo



You can install **SuMo** in two ways:



Option 1: Direct Installation via NPM

Run `npm install @morenabarboni/sumo`



Option 2: Add to package.json

And then run `npm install`.



`package/@morenabarboni/sumo`



`morenabarboni/sumo-solidity-mutator`

SuMo-Config.js

It allows developers to customize how Sumo should behave within a project.

```
module.exports = {  
  contractsDir: "auto",  
  testDir: "auto",  
  skipContracts: ["libraries"],  
  skipTests: [],  
  
  testingFramework: "auto",  
  testingTimeOutInSec: 500,  
  
  minimalOperators: false,  
  randomSampling: false,  
  randomMutants: 100  
}
```

Project Structure

Automatically detect or override contract, build, and test directories, and blacklists;

Test Execution





Configures test framework and timeout.

Mutation Strategy

Controls the scope and selection of mutants through minimal rules and random sampling.

Supported Testing Frameworks

SuMo is a **stand-alone module** that connects to the existing **testenv** to compile and test smart contracts, supporting broader project compatibility.

	 HardHat	 Foundry	 Brownie	 Custom
TEST LANGUAGE	JavaScript, TypeScript (mocha)	Solidity (forge)	Python (pytest)	Define Custom Test Script
CHAIN SIMULATOR	HardHat Network	Anvil	Ganache	-

Choosing Mutation Operators

Category	ID	Mutation Example
Types, Units, and Locations	AVR, DLR, VUR	<code>someAddress</code> → <code>address(0)</code>
Function Modifiers	MOD, MOI, OMD, PKD	<code>function pay()</code> <code>payable</code> → <code>function pay()</code>
Global Variables and Functions	GVR, TOR	<code>tx.origin</code> → <code>msg.sender</code>
Return Semantics	RSD, RVS	<code>return transfer();</code> → <code>transfer(); return true;</code>
Math, Crypto and Libraries	MCR, SFR	<code>safeMath.add()</code> → <code>safeMath.sub()</code>

Operator Selection (Individual, Cluster)

- Useful for targeting specific aspects of smart contract behavior;
- e.g., arithmetic logic, state visibility, control flow, events.

Generating Mutants

You can generate mutants without running tests: `npx sumo lookup`.

1. Parses Solidity source code into AST;
2. Explores nodes with each **operator** (custom visitor);
3. Applies **rule-based mutations** at matched nodes;
4. Generates a **mutations.json** file with all the mutants;

Contracts Summary



Contract	Total Mutants	Killed	Live	Stillborn	Timed Out	Untested
CampusCoin.sol	68	0	0	0	0	68

Mutant Pruning

Focus testing efforts, manage runtime and analysis cost.

1) Minimal Operators

- Each operator injects one mutation per target;
- Rules empirically found to be effective.

2) Random Mutant Selection

- Select a random subset of n mutants to be tested;

3) Coverage-Based Mutant Selection

- Target critical statements (with highest coverage);
- Target under-tested statements (with lowest coverage).

Running Mutation Testing

You can start mutation testing on the contracts under test by running `npx sumo test`. This actually starts the testing process:

Starting Mutation Testing

```
> Mutation 1 of 54 - [m7dff00e7 of CampusCoin.sol]
```

```
Applying mutation m7dff00e7 to CampusCoin.sol
```

```
119 |         totalSpent[msg.sender] += amount;  
    |         totalSpent[msg.sender] = amount;
```

```
Compiling mutation m7dff00e7 of CampusCoin.sol
```

```
npx hardhat compile
```

```
Compiled 1 Solidity file successfully (evm target: paris).
```

```
Running tests for mutant m7dff00e7
```

```
npx hardhat test --bail
```

- ☐ Find the complete results in: [SuMo-Demo/sumo-results](#)

What can SuMo tell us about CampusCoin?

The Mutation Score is average despite achieving full coverage ...

Mutation Score: 76.2%

Total Contracts: 1 | Total Mutants: 67

We can now **analyze live mutants** to derive **new tests** and improve the **fault-detection** of the CampusCoin.js test suite.

Mutation Analysis – Insight 1

All EED (Event-Emission-Deletion) mutants **survived** mutation testing.

```
function mint(address to, uint256 amount) public onlyAdmin {  
    require(isStudent[to], "Can only mint to registered students");  
    _mint(to, amount);  
    --- emit TokensMinted(to, amount);  
    +++ // emit TokensMinted(to, amount);  
}
```

EED mutant m95fd4816

Implication: Our tests never check the correct emission of events!

- **Events** → How contracts communicate with the outside world;
- **Faulty Events** → Off-chain systems may misinterpret key actions.

What about the Coverage Report?

In the Coverage Report all the **event emission statements** were covered at least once, giving us a false sense of security about their correctness.

```
function removeStudent(address student) external onlyAdmin {  
  1×    isStudent[student] = false;  
  1×    emit StudentRemoved(student);  
}
```

```
function removeServiceProvider(address provider) external onlyAdmin {  
  2×    serviceProviders[provider].active = false;  
  2×    emit ServiceProviderRemoved(provider);  
}
```

Updated Test to Kill EED Mutant

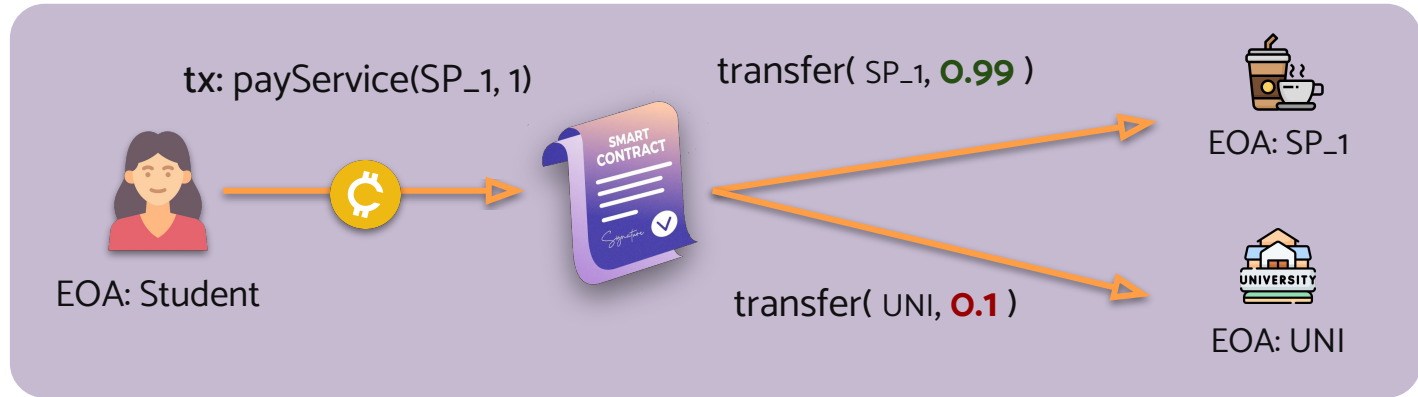
We used dedicated matchers to check for event emissions:

```
it("Should mint tokens to a student", async () {  
  await expect(campusCoin.mint(student1.address, "100"))  
    .to.emit(campusCoin, "TokensMinted")  
    .withArgs(student1.address, "100");  
  
  const balance = await campusCoin.balanceOf(student.address);  
  expect(balance).to.equal(ethers.parseUnits("100", 18));  
});
```

Test this with `npx sumo disable` ☐ `npx sumo enable EED` ☐ `npx sumo test`

☐ Find the enhanced test suite in: [solutions/3\)Test-EED-Mutants-Killer.js](#)

PayService() - Expected Behavior



A Student sends a payment to a Service Provider:

- 99% goes to the Service Provider;
- 1% fee is transferred to the University.

SuMo Report – Insight 2

Multiple **live mutants** around **fee-related logic** (both *fee computation* and *transfer*) suggest that this entire area of the contract is **under-tested**.

```
function payService(address to, uint256 amount) external {  
    ---| _transfer(msg.sender, university, fee);  
    +++| /* _transfer(msg.sender, university, fee); */  
}
```

FCD mutant m88a18980

- This drastic mutation removes the fee transfer entirely!
- Yet, none of our tests detect the issue ...
- Let's check the test method: "it: Should pay service"

SuMo Report – Insight 2

The enhanced test case exposed a bug in the original Smart Contract!

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount / 100) * UNIT; // fee computation  
}
```

Precision Loss Bug: Amount first divided by 100 then multiplied by UNIT:

- Test Input: amount = 1 token;
- Expected Fee: $1 / 100 = \mathbf{0.01 \text{ tokens}}$ //cannot represent with uint256
- Actual Fee: fee = 0 tokens //drops the fractional part

SuMo Report – Insight 2

The enhanced test case exposed a bug in the original Smart Contract!

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount / 100) * UNIT; // 1%  
}
```

BUGGY CONTRACT

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount * UNIT) / 100; // 1%  
}
```

FIXED CONTRACT

Conclusions

Mutation testing is a powerful complement to coverage analysis - it doesn't tell you *where your tests go*, but *what they actually prove*.

- Gives a more meaningful measure of test suite quality;
- Reduce the risk of undetected bugs before deployment.

There's no such thing as a free lunch:

- Running Mutation Testing is time consuming;
- Analyzing mutants can be overwhelming.

Mutant-Driven Test Generation Via LLMs

Alchemist automatically improves the **quality** of existing unit tests:

- Identifies **quality gaps** using **Mutation Testing**;
- Automatically fills these **quality gaps** using **LLMs**.

Alchemist: LLM-Driven Test Generation using Solidity Mutants and the Scientific Method

Morena Barboni*, Filippo Lampa*, Andrea Morichetta*, Andrea Polini*, and Edward Zulkoski†

*University of Camerino, {morena.barboni, filippo.lampa, andrea.morichetta, andrea.polini}@unicam.it

†Quantstamp, {ed}@quantstamp.com

Abstract—Bugs in Solidity smart contracts have led to significant financial losses, highlighting the importance of rigorous testing. Mutation testing is a powerful technique for evaluating test suite adequacy by identifying undetected faults introduced through small code changes. However, writing test cases for detecting live mutants is a labor-intensive task. This is especially true in the context of smart contracts, which involve complex interactions, access control considerations, and blockchain-specific behavior. To address this challenge, we propose Alchemist, a framework for generating Solidity test cases using Large Language Models (LLMs). Alchemist embeds the principles of the scientific method into the code generation process. This workflow can support the creation of more focused and interpretable mutant-killing tests, ultimately reducing developer effort.

Index Terms—Mutation Testing, Large Language Model, Test Generation, Ethereum, Smart Contract, Solidity

II. BACKGROUND

a) Mutation Testing in Smart Contracts: Mutation testing evaluates the adequacy of a test suite by introducing faults (i.e., *mutants*) into the code and checking whether existing tests can detect them. If a test case detects a mutant (i.e., the test fails), the mutant is considered *killed*. If a mutant goes undetected, it remains *live*, signaling a gap in the test suite. This approach goes beyond simple code coverage metrics [4] by revealing whether tests can capture behavioral deviations in the logic. Although effective, mutation testing is labor-intensive, especially in Solidity, where tools like Vertigo [5], SuMo [6], and ContractMut [7] generate domain-specific mutants that require careful analysis. Writing tests to kill these mutants remains a major barrier to wider adoption.