



UNIVERSITÀ
di CAMERINO

SuMo

Mutation Testing for Solidity Contracts



Test Adequacy Assessment

1. Introduction to Solidity SUT:

- Goal: Is the test suite good enough?

2. Code Coverage Analysis:

- What it tells us about our tests;
- Why it doesn't tell the whole story.

3. Mutation Testing with SuMo:

- Introduction to SuMo;
- Running Mutation Testing;
- Analyzing Live Mutants.



Demo Repository



morenabarboni/sumo-demo

SuMo-Demo

In the repository you will find:

- CampusCoin Setup
- Test Code Examples
- Slides

Prerequisites

Ensure the following are installed:

- Node.js
- npm (comes with Node.js)

SUT: Business Logic

contracts/CampusCoin.sol: A custom **ERC-20 token** that can be used to pay for services around campus (e.g., buying food, borrowing books).

Active Roles:

- **Admin:** Registers Users and Mints new CC;
- **Student:** Pays Service Providers using CC;

Passive Roles:

- **ServiceProvider:** Receives CC payments;
- **University:** Receives a 1% fee on payments.



The contract implements financial logic and must be carefully tested.

SUT: Test File

test/CampusCoin.js: verifies the core functionalities of CampusCoin, ensuring that it behaves as expected under typical usage scenarios.

👉 **Let's try it out: `npx hardhat test`**

Initial Impression: The test suite appears **relatively mature**:

- Covers major functionalities: deployment, transfers, payments ...
- Demonstrates awareness of the financial aspects of the contract.

SUT: Test File

test/CampusCoin.js: verifies the core functionalities of CampusCoin, ensuring that it behaves as expected under typical usage scenarios.

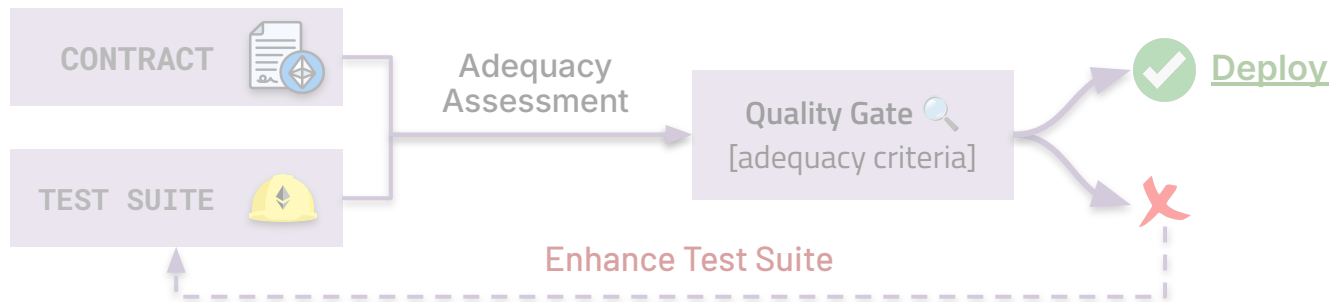
👉 **Let's try it out: `npx hardhat test`**

Initial Impression: The test suite appears **relatively mature**:

- Covers major functionalities: deployment, transfers, payments ...
- Demonstrates awareness of the financial aspects of the contract.

Goal

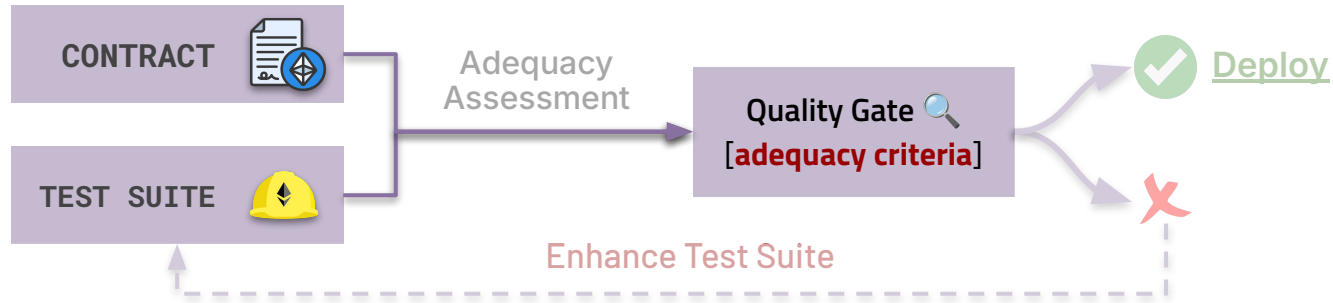
Does the CampusCoin.js test suite gives us **sufficient confidence** in the **correctness** of the CampusCoin.sol Smart Contract?



- To answer, we must establish what is “sufficient confidence”;
- That is, we must agree on some specific adequacy criteria.

Goal

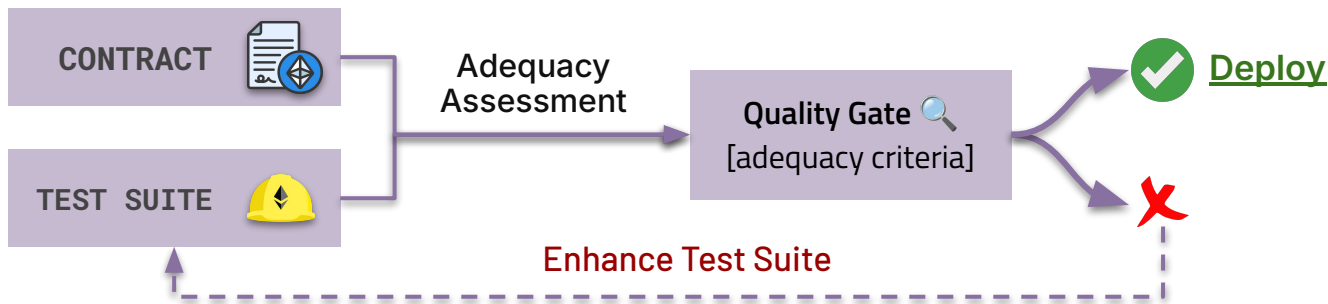
Does the CampusCoin.js test suite gives us **sufficient confidence** in the **correctness** of the CampusCoin.sol Smart Contract?



- To answer, we must establish what is “sufficient confidence”;
- That is, we must agree on some specific adequacy criteria.

Goal

Does the CampusCoin.js test suite gives us **sufficient confidence** in the **correctness** of the CampusCoin.sol Smart Contract?



- Then, we can run adequacy assessment on the test suite;
- If the criteria are met, the Contract passes the Quality Gate;

Quality Gate Based on Coverage Criteria

To ensure **high confidence** in the **correctness** of our contract, we define **strict coverage thresholds** for our deployment quality gate:



Statement Coverage == 100%

Every statement in the contract must be executed at least once.



Branch Coverage == 100%

Every branch (e.g., **require**) must be evaluated in both directions.

Adequacy Assessment - Coverage Analysis


solidity-coverage

chat on [gitter](#) npm@latest v0.8.16  FAILED  codecov 97%  Hardhat  Plugin

Code coverage for Solidity testing

```
22 modifier onlyColonyOwners {  
23   17× if (!this.userIsInRole(msg.sender, 0)) { throw; }  
24   17×  
25 }
```

- For more details about what this is, how it works and potential limitations, see [the accompanying article](#).
- `solidity-coverage` is [Solcover](#)



👉 Let's try it out: `npx hardhat coverage`

Coverage Analysis: Results

Looks like we're done, our test suite meets the **adequacy criteria**!



File	% Stmts	% Branch	% Funcs	% Lines
contracts\ CampusCoin.sol	100 100	100 100	100 100	100 100
All files	100	100	100	100

... Where is the catch?

Really stupid "smart contract" bug let hackers steal \$31 million ...

'Accidental' bug froze \$280 million worth of ether in Parity wallet
8 nov 2017 — Millions of dollars' worth of ether could be frozen on **Parity's** cryptocurrency **wallet** because one individual "accidentally" triggered a **bug**.

Coverage Analysis: Why it's not enough

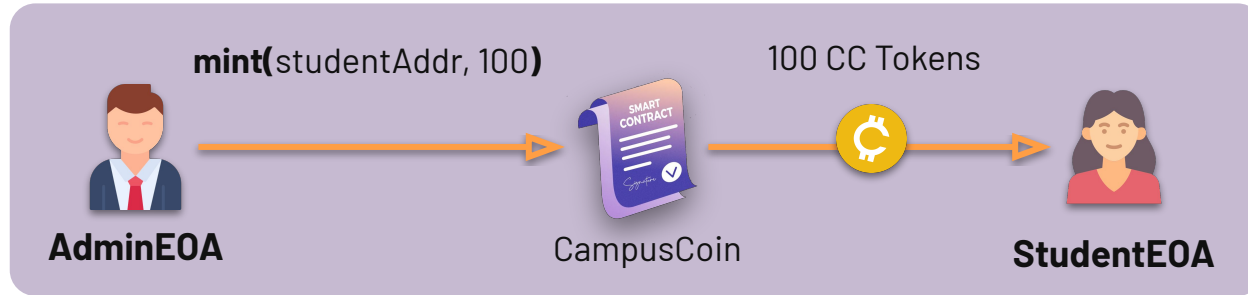
🛑 Most Developers would stop here !

File	% Stmts	% Branch	% Funcs	% Lines	Uncovered Lines
contracts\ CampusCoin.sol	100 100	100 100	100 100	100 100	
All files	100	100	100	100	

High coverage gives us a **false sense of security**:

- It only shows that certain parts of the code were executed;
- Not that they were meaningfully tested or properly verified.

Example: “Should mint tokens to a student”



CampusCoin implements a simple `mint(address, amount)` function:

- The Admin **mints** new CC tokens to a Student.
- The Student's **token balance** is updated accordingly.

Example: Test Case for mint()

```
it("Should mint tokens to student", async function () {  
    await campusCoin.mint(student1.address, "100");  
    const studentBalance = await campusCoin.balanceOf(student1.address);  
    expect(studentBalance).to.equal("100");  
});
```

CampusCoin.js: A **test method** verifies the correct behavior of `mint()`.

1. **Simulate Tx:** Admin mints 100 new tokens to a Student;
2. **Assert:** Confirm the correctness of the program behavior:
 - Actual output: new student balance
 - Expected output: 100 tokens

Example: "Should mint tokens to a student"

```
it("Should mint tokens to student", async function () {  
    await campusCoin.mint(student1.address, "100");  
    const studentBalance = await campusCoin.balanceOf(student1.address);  
    //expect(studentBalance).to.equal("100");  
});
```

What happens if we remove the assertion?

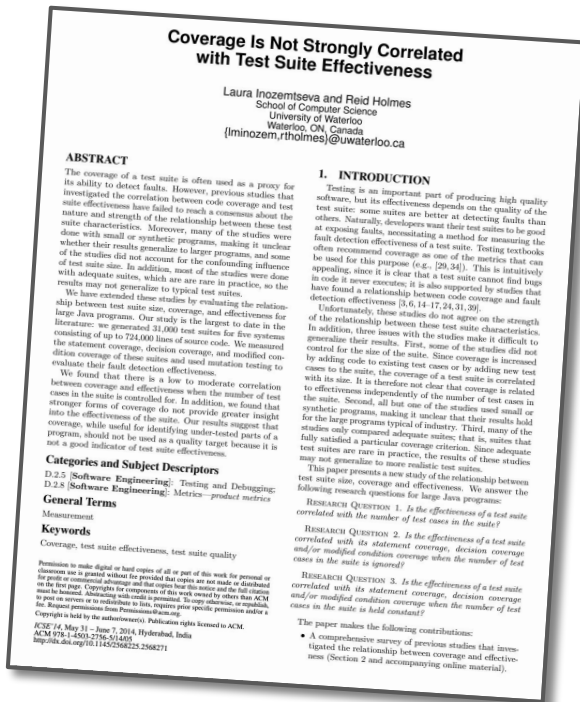
- We expect test adequacy to decrease!
- We are no longer enforcing any expectation;

👉 Let's try it out: `npx hardhat coverage`

Is Test Coverage Useless?

- **Good** for identifying under-tested parts of the system.
- **Bad** if used as a quality target!


— Inozemtseva and Holmes (2014)



Mutation Testing


Mutation testing offers a stronger alternative, as it evaluates a test suite based on its ability to detect *small faults* in the code.

- Tells us whether **assertions** are meaningful;
- If a **mutant survives**, the test suite **lacks a specific check** that can detect the fault;
- Encourages meaningful, bug-revealing tests.



SuMo - SOLidity MUtator

A **domain-aware** mutation testing framework that models **Solidity specific faults** to guide the derivation of more meaningful test cases.



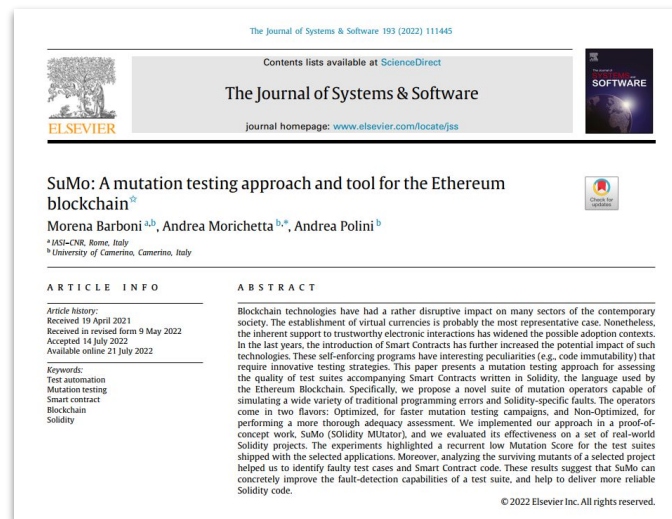
Automated Assessment

For any Solidity project
regardless of frameworks.



40 Mutation Operators


SOTA tools (e.g., PIT);
Solidity-centric analysis;
Continuous refinement.



Empirical Evaluation - Key Insights

Solidity Tests Show Large Oracle Gaps¹

Average Cov. = ~84%, Average MS = ~64%



- **Implication:**
Tests execute the code but don't always verify correctness.
- **Mind the Gap!**
It acts as a **diagnostic**, showing where tests fail to verify behavior, guiding testing efforts.

¹ Mind the gap: The difference between coverage and mutation score can guide testing efforts."

Empirical Evaluation - Key Insights

The Gap is larger for Solidity-Specific Mutants:

Traditional (MS = 68,2 %) | Solidity (MS = 61,2 %)


Target	MS	Potential Impact
Event	34%	Monitoring of contract behavior
Modifiers	37,7%	Access Control and reusable logic
Exception Handling	40,1%	Management of critical transaction reverting scenarios.
Blockchain Variable	64,5%	Any logic dependant on global blockchain properties (e.g., time)

Implications:

- Testers overlook Solidity constructs, as they introduce complex/unique semantics.
- Domain-relevant faults may remain untested and go unnoticed until deployment.

Mutation Testing can highlight these critical gaps that would otherwise be ignored.

What can SuMo tell us about CampusCoin?



You can install **SuMo** via NPM:

1. In **SuMo-Demo**, run **npm install**;
2. This will create **sumo-config.js**;



package/@morenabarboni/sumo



morenabarboni/sumo-solidity-mutator

SuMo-Config.js

It allows developers to customize how Sumo should behave within a project.

```
module.exports = {  
  contractsDir: "auto",  
  testDir: "auto",  
  skipContracts: ["libraries"],  
  skipTests: [],  
  
  testingFramework: "auto",  
  testingTimeOutInSec: 500,  
  
  minimalOperators: false,  
  randomSampling: false,  
  randomMutants: 100  
}
```

Project Structure

Automatically detect or override contract, build, and test directories, and blacklists;

Test Execution





Configures test framework and timeout.

Mutation Strategy

Controls the scope and selection of mutants through minimal rules and random sampling.

Supported Testing Frameworks

SuMo is a **stand-alone module** that connects to the existing **testenv** to compile and test smart contracts, supporting broader project compatibility.

	 HardHat	 Foundry	 Brownie	 Custom
TEST LANGUAGE	JavaScript, TypeScript (mocha)	Solidity (forge)	Python (pytest)	Define Custom Test Script
CHAIN SIMULATOR	HardHat Network	Anvil	Ganache	-

Choosing Mutation Operators

Category	ID	Mutation Example
Types, Units, and Locations	AVR, DLR, VUR	<code>someAddress</code> → <code>address(0)</code>
Function Modifiers	MOD, MOI, OMD, PKD	<code>function pay()</code> <code>payable</code> → <code>function pay()</code>
Global Variables and Functions	GVR, TOR	<code>tx.origin</code> → <code>msg.sender</code>
Return Semantics	RSD, RVS	<code>return transfer();</code> → <code>transfer(); return true;</code>
Math, Crypto and Libraries	MCR, SFR	<code>safeMath.add()</code> → <code>safeMath.sub()</code>

Operator Selection (Individual, Cluster)

- Useful for targeting specific aspects of smart contract behavior;
- e.g., arithmetic logic, state visibility, control flow, events.

Generating Mutants

You can generate mutants without running tests: `npx sumo lookup`.

1. Parses Solidity source code into AST;
2. Explores nodes with each **operator** (custom visitor);
3. Applies **rule-based mutations** at matched nodes;
4. Generates a **mutations.json** file with all the mutants;

Contracts Summary



Contract	Total Mutants	Killed	Live	Stillborn	Timed Out	Untested
CampusCoin.sol	68	0	0	0	0	68

Running Mutation Testing

You can start mutation testing on the contracts under test by running `npx sumo test`. This actually starts the testing process:

Starting Mutation Testing

```
> Mutation 1 of 54 - [m7dff00e7 of CampusCoin.sol]
```

```
Applying mutation m7dff00e7 to CampusCoin.sol
```

```
119 |         totalSpent[msg.sender] += amount;  
    |         totalSpent[msg.sender] = amount;
```

```
Compiling mutation m7dff00e7 of CampusCoin.sol
```

```
npx hardhat compile
```

```
Compiled 1 Solidity file successfully (evm target: paris).
```

```
Running tests for mutant m7dff00e7
```

```
npx hardhat test --bail
```

- ☐ Find the complete results in: [SuMo-Demo/sumo-results](#)

What can SuMo tell us about CampusCoin?

The MS is average despite achieving 100% coverage ...

Mutation Score: 76.2%

Total Contracts: 1 | Total Mutants: 67

Mutant Analysis: We now **analyze live mutants** and derive **new tests** to improve **fault-detection** in our test suite.

- ☐ Find the report in: [SuMo-Demo/sumo-results/index.html](https://sumo-demo.univcam.it/sumo-results/index.html)

Mutation Analysis – Insight 1

All EED (Event-Emission-Deletion) mutants **survived** mutation testing.

```
function mint(address to, uint256 amount) public onlyAdmin {  
    require(isStudent[to], "Can only mint to registered students");  
    _mint(to, amount);  
    --- emit TokensMinted(to, amount);  
    +++ // emit TokensMinted(to, amount);  
}
```

EED mutant m95fd4816

Implication: Our tests never check the correct emission of events!

- **Events** → How contracts communicate with the outside world;
- **Faulty Events** → Off-chain systems may misinterpret key actions.

In the Coverage Report all the **events** were covered, but **no test** asserts their **emission** or whether they **log expected arguments**.

```
function removeServiceProvider(address provider) external onlyAdmin {
2x   serviceProviders[provider].active = false;
2x   emit ServiceProviderRemoved(provider);
}
```

- Find this in: [SuMo-Demo/coverage/index.html](#)

Updated Test to Kill EED Mutant

We add assertions to check for missing event emissions:

Mutation Testing completed in 26 seconds 🖱️

SuMo generated 6 mutants:

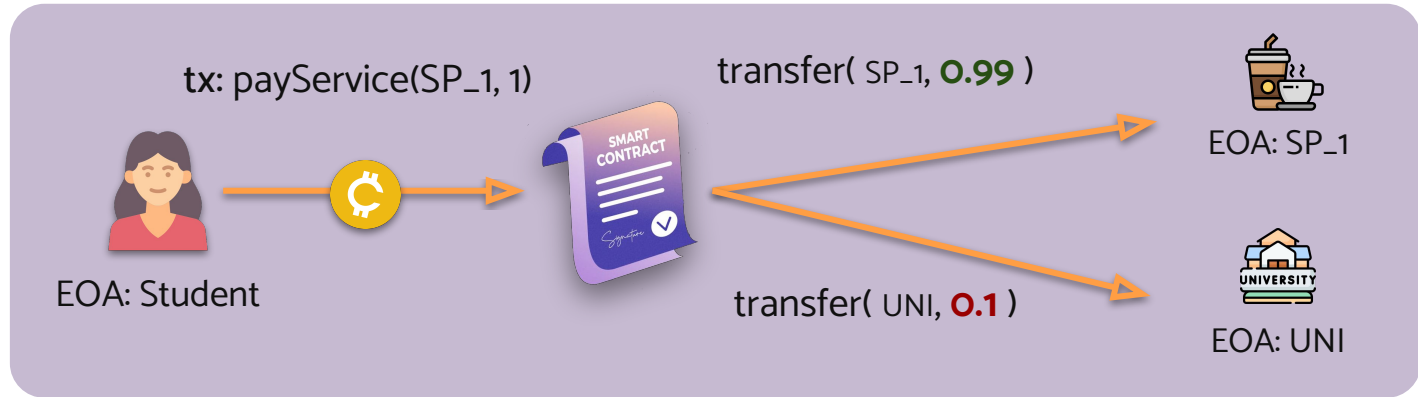
- 0 live;
- 6 killed;
- 0 stillborn;
- 0 timed-out.

Mutation Score: **100.00 %**

You can test this with:

- `npx sumo disable`
- `npx sumo enable EED`
- `npx sumo test`

PayService() - Expected Behavior



A Student sends a payment to a Service Provider:

- 99% goes to the Service Provider;
- 1% fee is transferred to the University.

SuMo Report – Insight 2

Multiple **live mutants** around **fee-related logic** (both *fee computation* and *transfer*) suggest that this entire area of the contract is **under-tested**.

```
function payService(address to, uint256 amount) external {  
    ---| _transfer(msg.sender, university, fee);  
    +++| /* _transfer(msg.sender, university, fee); */  
}
```

FCD mutant m88a18980

- This drastic mutation removes the fee transfer entirely!
- Yet, none of our tests detect the issue ...
- Let's check the test method: "it: Should pay service"

SuMo Report – Insight 2

The enhanced test case exposed a bug in the original Smart Contract!

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount / 100) * UNIT; // fee computation  
}
```

Precision Loss Bug: Amount first divided by 100 then multiplied by UNIT:

- Test Input: amount = 1 token;
- Expected Fee: $1 / 100 = \mathbf{0.01 \text{ tokens}}$ //cannot represent with uint256
- Actual Fee: fee = 0 tokens //drops the fractional part

SuMo Report – Insight 2

The enhanced test case exposed a bug in the original Smart Contract!

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount / 100) * UNIT; // 1%  
}
```

BUGGY CONTRACT

```
function payService(address to, uint256 amount) external {  
    uint256 fee = (amount * UNIT) / 100; // 1%  
}
```

FIXED CONTRACT

Conclusions

Mutation testing is a powerful complement to coverage analysis - it doesn't tell you *where your tests go*, but *what they actually prove*.

- Gives a more meaningful measure of test suite quality;
- Reduce the risk of undetected bugs before deployment.

There's no such thing as a free lunch:

- Running Mutation Testing is time consuming;
- Analyzing mutants can be overwhelming.